

# AN7914 Week 4 Python

March 1, 2023

## 1 Week 4 Python

### 1.1 Introduction to Pandas

**Pandas** is a package built on top of **NumPy**, and provides an efficient implementation of a **DataFrame**. **DataFrames** are essentially multidimensional *arrays* with attached row and column labels, and often with heterogeneous types and/or missing data. `### Installing Pandas` Installation of Pandas on your system requires NumPy to be installed. Details on this installation can be found in the [Pandas](#) documentation. Once Pandas is installed, you can import it and check the version:

```
[1]: import pandas
```

```
[2]: pandas.__version__
```

```
[2]: '1.5.2'
```

We will however use an alias to call pandas. So when importing we do the following

```
[3]: import pandas as pd
```

In the code above we imported **pandas** under the alias **pd**. Now let's check the version again, but this we will use the alias.

```
[4]: pd.__version__
```

```
[4]: '1.5.2'
```

### 1.2 Creating data

There are two core objects in pandas: the **DataFrame** and the **Series**. `### DataFrame` A **DataFrame** is a table. It contains an array of individual entries, each of which has a certain value. Each entry corresponds to a *row* (or record) and a *column*.

For example, consider the following simple DataFrame:

```
[5]: pd.DataFrame({'Yup': [50, 21,32], 'Nope': [131, 2,200]})
```

```
[5]:
```

	Yup	Nope
0	50	131
1	21	2

2 32 200

In this example, the “0, No” entry has the value of 131. The “0, Yes” entry has a value of 50, and so on.

`DataFrame` entries are not limited to integers. For instance, here’s a `DataFrame` whose values are strings:

```
[6]: pd.DataFrame({'Bob': ['I liked it.', 'It was awful.'], 'Sue': ['Pretty good.',  
↪ 'Bland.']}])
```

```
[6]:           Bob           Sue  
0    I liked it.  Pretty good.  
1  It was awful.    Bland.
```

We are using the `pd.DataFrame()` constructor to generate these `DataFrame` objects. The syntax for declaring a new one is a **dictionary** whose keys are the *column* names (Bob and Sue in this example), and whose *values* are a list of entries.

The dictionary-list constructor assigns values to the column labels, but just uses an ascending count from 0 (0, 1, 2, 3, ...) for the row labels. Sometimes this is OK, but oftentimes we will want to assign these labels ourselves.

```
[7]: pd.DataFrame({'Bob': ['I liked it.', 'It was awful.'],  
                  'Sue': ['Pretty good.', 'Bland.'],  
                  index=['Product A', 'Product B'])
```

```
[7]:           Bob           Sue  
Product A    I liked it.  Pretty good.  
Product B  It was awful.    Bland.
```

### 1.2.1 Series

A **Series**, by contrast, is a sequence of data values. If a `DataFrame` is a table, a `Series` is a **list**. And in fact you can create one with nothing more than a list:

```
[8]: pd.Series([1, 2, 3, 4, 5])
```

```
[8]: 0    1  
    1    2  
    2    3  
    3    4  
    4    5  
dtype: int64
```

A **Series** is, in essence, a single **column** of a `DataFrame`. So you can assign row labels to the `Series` the same way as before, using an `index` parameter. However, a **Series** does not have a column name, it only has one overall **name**:

```
[9]: pd.Series([30, 35, 40], index=['2015 Sales', '2016 Sales', '2017 Sales'],
      ↪name='Product A')
```

```
[9]: 2015 Sales    30
      2016 Sales    35
      2017 Sales    40
      Name: Product A, dtype: int64
```

The `Series` and the `DataFrame` are intimately related. It's helpful to think of a `DataFrame` as actually being just a bunch of `Series` “glued together”.

### 1.3 Importing data sets

You will need to use and import data sets from the internet or from your hard-drive. So if you want to import a `csv` file you will need to use `pd.read_csv()` command. The argument in the command could be the location where the file is stored in your computer or it could be a file store in the internet as show below.

```
[10]: df_tips = pd.read_csv('https://raw.githubusercontent.com/mwaskom/seaborn-data/
      ↪master/tips.csv')
```

We use `pd.read_csv()` to read a file stored in `'https://raw.githubusercontent.com/mwaskom/seaborn-data/m'`. Then we store this dataset as a `dataframe` in `df_tips`.

Now lets take a look at the dataset. We can simply type `df_tips` the name of the `dataframe`. It is not going to show all columns and rows.

```
[11]: df_tips
```

```
[11]:   total_bill  tip  sex smoker  day  time  size
0      16.99  1.01 Female    No  Sun  Dinner    2
1      10.34  1.66  Male    No  Sun  Dinner    3
2      21.01  3.50  Male    No  Sun  Dinner    3
3      23.68  3.31  Male    No  Sun  Dinner    2
4      24.59  3.61 Female    No  Sun  Dinner    4
..      ...  ...  ...    ...  ...  ...  ...
239     29.03  5.92  Male    No  Sat  Dinner    3
240     27.18  2.00 Female   Yes  Sat  Dinner    2
241     22.67  2.00  Male   Yes  Sat  Dinner    2
242     17.82  1.75  Male    No  Sat  Dinner    2
243     18.78  3.00 Female    No  Thur Dinner    2
```

```
[244 rows x 7 columns]
```

If you want to know exactly how many rows and columns the `dataframe` has we can simply type `df_tips.shape`

```
[12]: df_tips.shape
```

```
[12]: (244, 7)
```

We see that the output is (244, 7)– this means we have 244 rows and 7 columns.

Now let's take a look at the first 15 rows.

```
[13]: df_tips.head(15)
```

```
[13]:
```

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4
5	25.29	4.71	Male	No	Sun	Dinner	4
6	8.77	2.00	Male	No	Sun	Dinner	2
7	26.88	3.12	Male	No	Sun	Dinner	4
8	15.04	1.96	Male	No	Sun	Dinner	2
9	14.78	3.23	Male	No	Sun	Dinner	2
10	10.27	1.71	Male	No	Sun	Dinner	2
11	35.26	5.00	Female	No	Sun	Dinner	4
12	15.42	1.57	Male	No	Sun	Dinner	2
13	18.43	3.00	Male	No	Sun	Dinner	4
14	14.83	3.02	Female	No	Sun	Dinner	2

`df_tips.head(15)` gives us the first 15 rows of the dataframe. If we typed `df_tips.head(25)` it would show us the first 25 rows.

To see the last 15 rows we can use `df_tips.tail(15)`

```
[14]: df_tips.tail(15)
```

```
[14]:
```

	total_bill	tip	sex	smoker	day	time	size
229	22.12	2.88	Female	Yes	Sat	Dinner	2
230	24.01	2.00	Male	Yes	Sat	Dinner	4
231	15.69	3.00	Male	Yes	Sat	Dinner	3
232	11.61	3.39	Male	No	Sat	Dinner	2
233	10.77	1.47	Male	No	Sat	Dinner	2
234	15.53	3.00	Male	Yes	Sat	Dinner	2
235	10.07	1.25	Male	No	Sat	Dinner	2
236	12.60	1.00	Male	Yes	Sat	Dinner	2
237	32.83	1.17	Male	Yes	Sat	Dinner	2
238	35.83	4.67	Female	No	Sat	Dinner	3
239	29.03	5.92	Male	No	Sat	Dinner	3
240	27.18	2.00	Female	Yes	Sat	Dinner	2
241	22.67	2.00	Male	Yes	Sat	Dinner	2
242	17.82	1.75	Male	No	Sat	Dinner	2
243	18.78	3.00	Female	No	Thur	Dinner	2

If we want to see the names of the columns we use `df_tips.columns`

```
[15]: df_tips.columns
```

```
[15]: Index(['total_bill', 'tip', 'sex', 'smoker', 'day', 'time', 'size'],  
dtype='object')
```

If we want to see a specific columns we can pass in a list – `df_tips[['total_bill','tip']]`. We passed in the list `['total_bill','tip']`. This list contains the list of column names. The output is going to look a `dataframe` and not a `series`.

```
[16]: df_tips[['total_bill','tip']]
```

```
[16]:
```

	total_bill	tip
0	16.99	1.01
1	10.34	1.66
2	21.01	3.50
3	23.68	3.31
4	24.59	3.61
..	...	...
239	29.03	5.92
240	27.18	2.00
241	22.67	2.00
242	17.82	1.75
243	18.78	3.00

[244 rows x 2 columns]

```
[ ]:
```